

AI Safety & Red Teaming — Security Report

Target: Project 01 — conversational e-commerce assistant (deterministic fallback, no LLM)

Generated 2026-09-29 02:18 UTC · OWASP LLM Top 10 · corpus of 106 attacks

Live HTTP run against the real system started locally, but the target had no LLM configured (llm_enabled=false): it answered with its deterministic fallback. This measures the framework and the guardrail reduction, not the target's quality under a real model.

Executive summary

Metric	Value
Attack success rate (baseline)	0.0%
Attack success rate (with guardrails)	0.0%
Absolute reduction	0.0%
Guardrail false-positive rate	0.0%
Mean guardrail latency	1.00 ms
Attacks executed	106
Legitimate queries tested	100

Attack success by OWASP category

OWASP category	Baseline	Guardrails	FP rate (domain)
LLM01 Prompt Injection	0.0%	0.0%	0.0%
LLM02 Insecure Output Handling	0.0%	0.0%	0.0%
LLM03 Training Data Poisoning	0.0%	0.0%	0.0%
LLM04 Model Denial of Service	0.0%	0.0%	0.0%
LLM05 Supply Chain Vulnerabilities	0.0%	0.0%	0.0%
LLM06 Sensitive Information Disclosure	0.0%	0.0%	0.0%
LLM07 Insecure Plugin Design	0.0%	0.0%	0.0%
LLM08 Excessive Agency	0.0%	0.0%	0.0%
LLM09 Overreliance	0.0%	0.0%	0.0%
LLM10 Model Theft	0.0%	0.0%	0.0%

Top findings (successful attacks, by severity)

No attacks succeeded against the baseline target.

Methodology

Each attack is fired at the target, the response is analysed for refusal, compliance markers, PII leakage and denial-of-service, and mapped to its OWASP category. Guardrails screen inputs (injection/jailbreak patterns) and sanitise outputs (PII redaction, HTML neutralisation). Ground truth for every attack is hand-written. False-positive rate is measured over a benign query set.

Run IDs: baseline=chat-baseline, guardrailed=chat-guardrailed, fp=chat-fp.